

Importing the Abalone Dataset

In [57]: `import pandas as pd`

In [3]: `url = ("https://archive.ics.uci.edu/ml/machine-learning-databases/abalone/abalone.csv")`

In [4]: `abalone = pd.read_csv(url, header=None)`

In [5]: `abalone.head()`

Out[5]:

	0	1	2	3	4	5	6	7	8
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7


In [6]: `abalone.columns = ["Sex", "Length", "Diameter", "Height", "Whole weight", "Shucked weight", "Viscera weight", "Shell weight", "Rings"]`

In [7]: `abalone = abalone.drop("Sex", axis=1)`

Descriptive Statistics From the Abalone Dataset

In [8]: `import matplotlib.pyplot as plt`

```
In [9]: abalone["Rings"].hist(bins=15)  
plt.show()
```



```
In [10]: correlation_matrix = abalone.corr()
```

```
In [11]: correlation_matrix["Rings"]
```

```
Out[11]: Length          0.556720  
Diameter        0.574660  
Height          0.557467  
Whole weight    0.540390  
Shucked weight  0.420884  
Viscera weight  0.503819  
Shell weight    0.627574  
Rings           1.000000  
Name: Rings, dtype: float64
```

A Step-by-Step kNN

```
In [12]: import numpy as np
```

```
In [13]: a = np.array([2, 2])
```

```
In [14]: b = np.array([4, 4])
```

```
In [15]: np.linalg.norm(a - b)
```

```
Out[15]: 2.8284271247461903
```

k Nearest Neighbors#

```
In [16]: X = abalone.drop("Rings", axis=1)
```

```
In [17]: X = X.values
```

```
In [18]: y = abalone["Rings"]
```

```
In [19]: y = y.values
```

```
In [20]: new_data_point = np.array([0.569552, 0.446407, 0.154437, 1.016849, 0.439051, 0.222526,
```

```
In [21]: distances = np.linalg.norm(X - new_data_point, axis=1)
```

```
In [22]: k = 3  
nearest_neighbor_ids = distances.argsort()[:k]  
nearest_neighbor_ids
```

```
Out[22]: array([4045, 1902, 1644], dtype=int64)
```

Voting or Averaging of Multiple Neighbors

```
In [23]: nearest_neighbor_rings = y[nearest_neighbor_ids]
nearest_neighbor_rings
```

```
Out[23]: array([ 9, 11, 10], dtype=int64)
```

Average for Regression

```
In [24]: prediction = nearest_neighbor_rings.mean()
```

Mode for Classification

```
In [25]: import scipy.stats
```

```
In [26]: class_neighbors = np.array(["A", "B", "B", "C"])
```

```
In [27]: scipy.stats.mode(class_neighbors)
```

C:\Users\Mona Shah\AppData\Local\Temp\ipykernel_14716\494549111.py:1: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
scipy.stats.mode(class_neighbors)
```

C:\Users\Mona Shah\AppData\Roaming\Python\Python310\site-packages\scipy\stats_stats_py.py:110: RuntimeWarning: The input array could not be properly checked for nan values. nan values will be ignored.

```
warnings.warn("The input array could not be properly "
```

C:\Users\Mona Shah\AppData\Local\Temp\ipykernel_14716\494549111.py:1: DeprecationWarning: Support for non-numeric arrays has been deprecated as of SciPy 1.9.0 and will be removed in 1.11.0. `pandas.DataFrame.mode` can be used instead, see <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.mode.html>. (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.mode.html>.)

```
scipy.stats.mode(class_neighbors)
```

```
Out[27]: ModeResult(mode=array(['B'], dtype='<U1'), count=array([2]))
```

kNN in Python Using scikit-learn

```
In [28]: from sklearn.model_selection import train_test_split
```

```
In [29]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s
```

kNN Regression in scikit-learn to the Abalone Dataset

```
In [30]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [31]: knn_model = KNeighborsRegressor(n_neighbors=3)
```

```
In [32]: knn_model.fit(X_train, y_train)
```

```
Out[32]: KNeighborsRegressor
KNeighborsRegressor(n_neighbors=3)
```

Using scikit-learn to Inspect Model Fit

```
In [33]: from sklearn.metrics import mean_squared_error
```

```
In [34]: from math import sqrt
```

```
In [35]: train_preds = knn_model.predict(X_train)
```

```
In [36]: mse = mean_squared_error(y_train, train_preds)
```

```
In [37]: rmse = sqrt(mse)
rmse
```


```
Out[37]: 1.6538366794859511
```

```
In [38]: test_preds = knn_model.predict(X_test)
mse = mean_squared_error(y_test, test_preds)
rmse = sqrt(mse)
rmse
```


```
Out[38]: 2.375417924000521
```

Plotting

```
In [39]: import seaborn as sns
cmap = sns.cubehelix_palette(as_cmap=True)
f, ax = plt.subplots()
points = ax.scatter(X_test[:, 0], X_test[:, 1], c=test_preds, s=50, cmap=cmap)
f.colorbar(points)
plt.show()
```



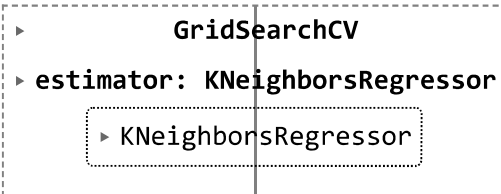
```
In [40]: cmap = sns.cubehelix_palette(as_cmap=True)
f, ax = plt.subplots()
points = ax.scatter(X_test[:, 0], X_test[:, 1], c=y_test, s=50, cmap=cmap)
f.colorbar(points)
plt.show()
```



Improving kNN Performances in scikit-learn Using GridSearchCV

```
In [43]: from sklearn.model_selection import GridSearchCV
parameters = {"n_neighbors": range(1, 50)}
gridsearch = GridSearchCV(KNeighborsRegressor(), parameters)
gridsearch.fit(X_train, y_train)
```

```
Out[43]:
```



```
In [44]: gridsearch.best_params_
```

```
Out[44]: {'n_neighbors': 25}
```

```
In [45]: train_preds_grid = gridsearch.predict(X_train)
train_mse = mean_squared_error(y_train, train_preds_grid)
train_rmse = sqrt(train_mse)
test_preds_grid = gridsearch.predict(X_test)
test_mse = mean_squared_error(y_test, test_preds_grid)
test_rmse = sqrt(test_mse)
train_rmse
```

```
Out[45]: 2.0731180327543384
```

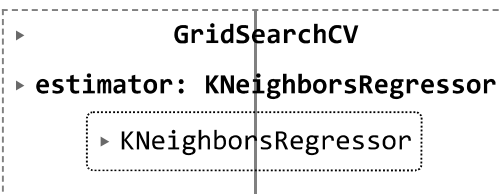
```
In [46]: test_rmse
```

```
Out[46]: 2.1700197339962175
```

Adding Weighted Average of Neighbors Based on Distance

```
In [47]: parameters = {"n_neighbors": range(1, 50),
"weights": ["uniform", "distance"],}
gridsearch = GridSearchCV(KNeighborsRegressor(), parameters)
gridsearch.fit(X_train, y_train)
```

```
Out[47]:
```



```
In [48]: gridsearch.best_params_
```

```
Out[48]: {'n_neighbors': 25, 'weights': 'distance'}
```



```
In [49]: test_preds_grid = gridsearch.predict(X_test)
test_mse = mean_squared_error(y_test, test_preds_grid)
test_rmse = sqrt(test_mse)
test_rmse
```

Out[49]: 2.1634265584947485

Improving on kNN in scikit-learn With Bagging

```
In [51]: best_k = gridsearch.best_params_["n_neighbors"]
best_weights = gridsearch.best_params_["weights"]
bagged_knn = KNeighborsRegressor(n_neighbors=best_k, weights=best_weights)
```

```
In [54]: from sklearn.ensemble import BaggingRegressor
bagging_model = BaggingRegressor(bagged_knn, n_estimators=100)
```

```
In [56]: test_mse = mean_squared_error(y_test, test_preds_grid)
test_rmse = sqrt(test_mse)
test_rmse
```

Out[56]: 2.1634265584947485

Conclusion

With the Above analysis we can say that using KNN model with bagging optimizes the performance to the maximum

Also with the above python libraries we can get the best nearest neighbour analysis with minimum code

In []: